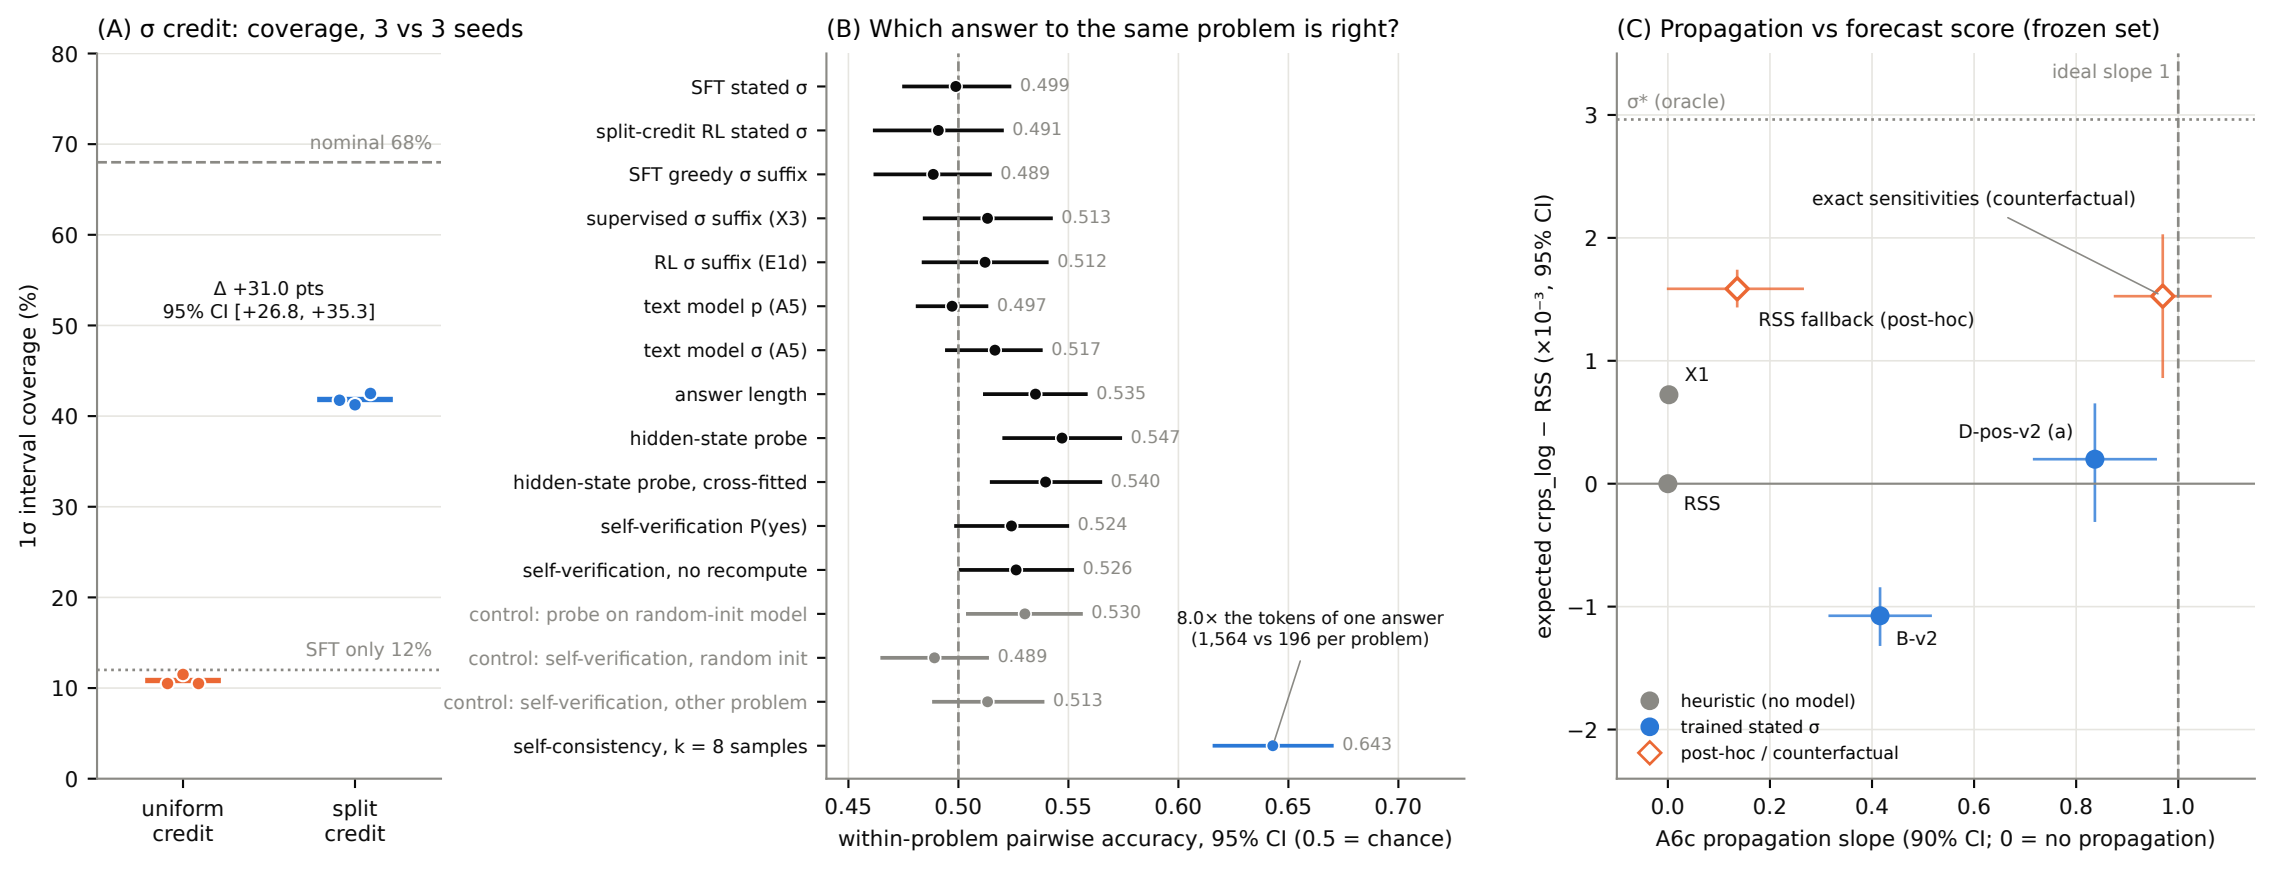
(A) σ credit: coverage, 3 vs 3 seeds
1σ interval coverage (%)
nominal 68%
Δ +31.0 pts
95% CI [+26.8, +35.3]
SFT only 12%
uniform credit
split credit
(B) Which answer to the same problem is right?
SFT stated σ
split-credit RL stated σ
SFT greedy σ suffix
supervised σ suffix (X3)
RL σ suffix (E1d)
text model p (A5)
text model σ (A5)
answer length
hidden-state probe
hidden-state probe, cross-fitted
self-verification P(yes)
self-verification, no recompute
control: probe on random-init model
control: self-verification, random init
control: self-verification, other problem
self-consistency, k = 8 samples
0.499
0.491
0.489
0.513
0.512
0.497
0.517
0.535
0.547
0.540
0.524
0.526
0.530
0.489
0.513
0.643
8.0× the tokens of one answer
(1,564 vs 196 per problem)
within-problem pairwise accuracy, 95% CI (0.5 = chance)
(C) Propagation vs forecast score (frozen set)
ideal slope 1
σ* (oracle)
exact sensitivities (counterfactual)
RSS fallback (post-hoc)
X1
D-pos-v2 (a)
RSS
B-v2
heuristic (no model)
trained stated σ
post-hoc / counterfactual
expected crps_log − RSS (×10⁻³, 95% CI)
A6c propagation slope (90% CI; 0 = no propagation)